


Measured GPU power-performance trade-offs

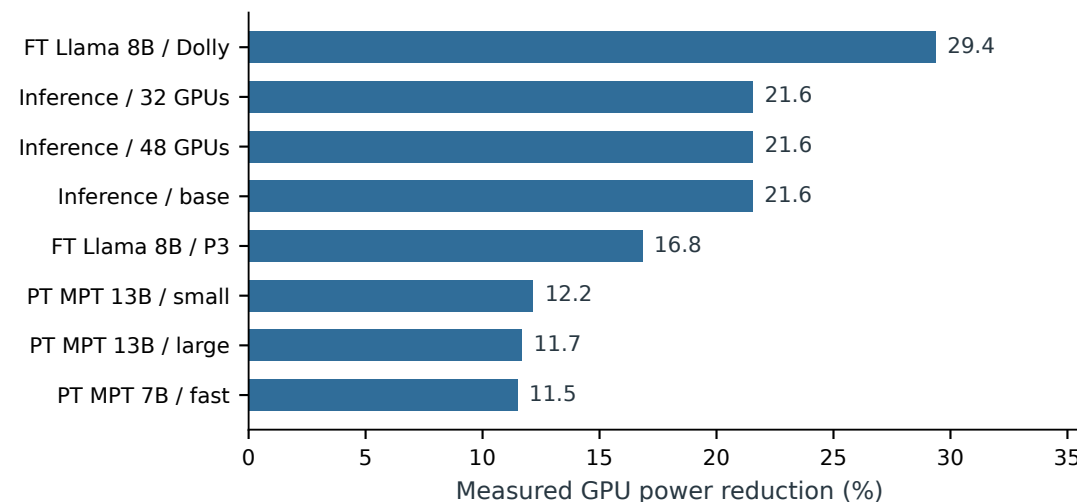
Published experiment data; independent calculations. No new hardware experiment or China-wide inference.

a Measured response curves

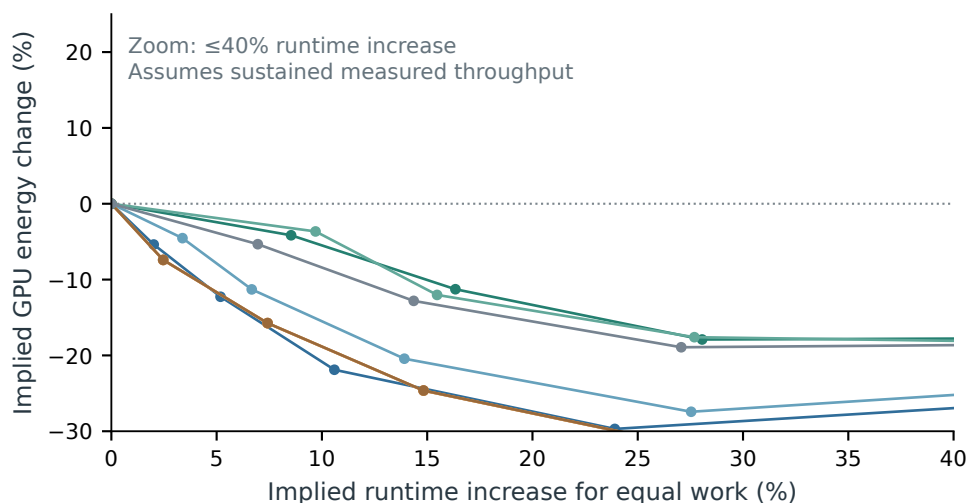


FT Llama 8B / Dolly
 FT Llama 8B / P3
 Inference / 48 GPUs
 Inference / 32 GPUs
 Inference / base
 PT MPT 13B / large
 PT MPT 13B / small
 PT MPT 7B / fast

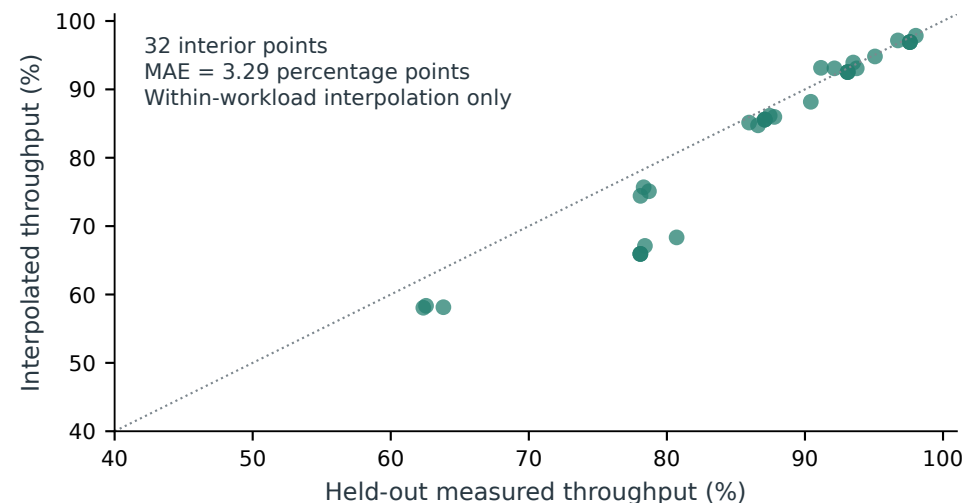
b Power reduction with throughput $\geq 90\%$



c Equal-work GPU energy implication



d Interior-point interpolation check



Source: Colangelo et al., Nature Energy; ai-emerald/emerald-ai-demo-may-2025, dvfs_sweep.csv.

90% is an analyst-selected throughput screen, not a verified SLA. Eight configurations are not eight independent model samples.