

Public production GPU-cluster traces: waits are a lower bound on tolerated delay, not deadlines; no power telemetry in any trace

